AI/ML Model Interpretability

Understanding AI Model Internals

- Goal: Make AI models transparent and interpretable through systematic analysis
- Challenge: Understanding AI models is similar to biological research - complex systems requiring sophisticated tools

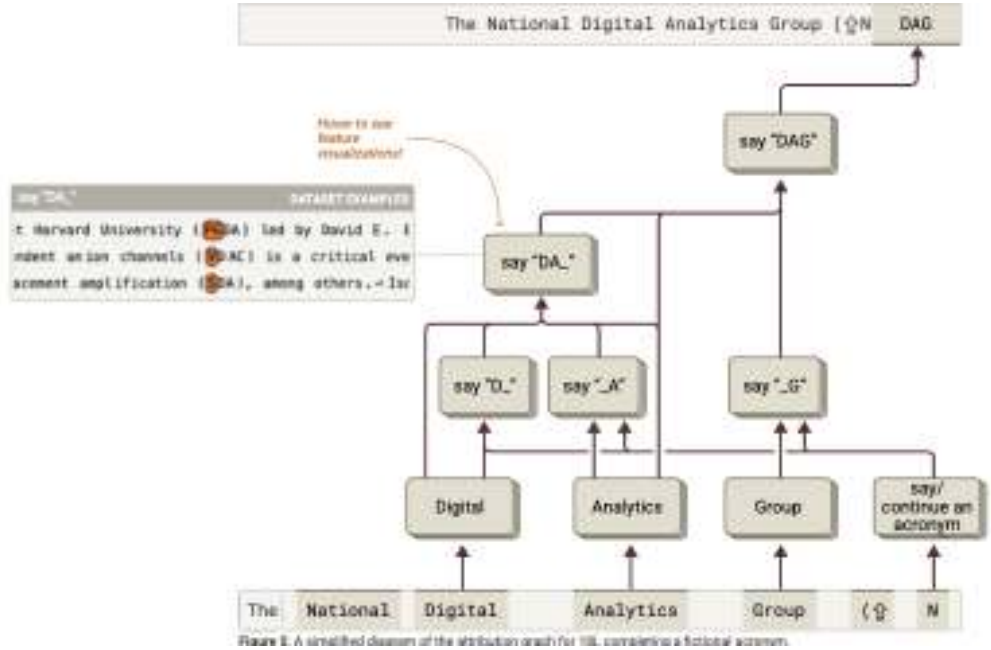

- Predictions about unexpected AI outputs
- "Microscopes" for AI model internals

Lindsey, et al., "On the Biology of a Large Language Model", Transformer Circuits, 2025.

One slide overview (MLP=ANN=FF)

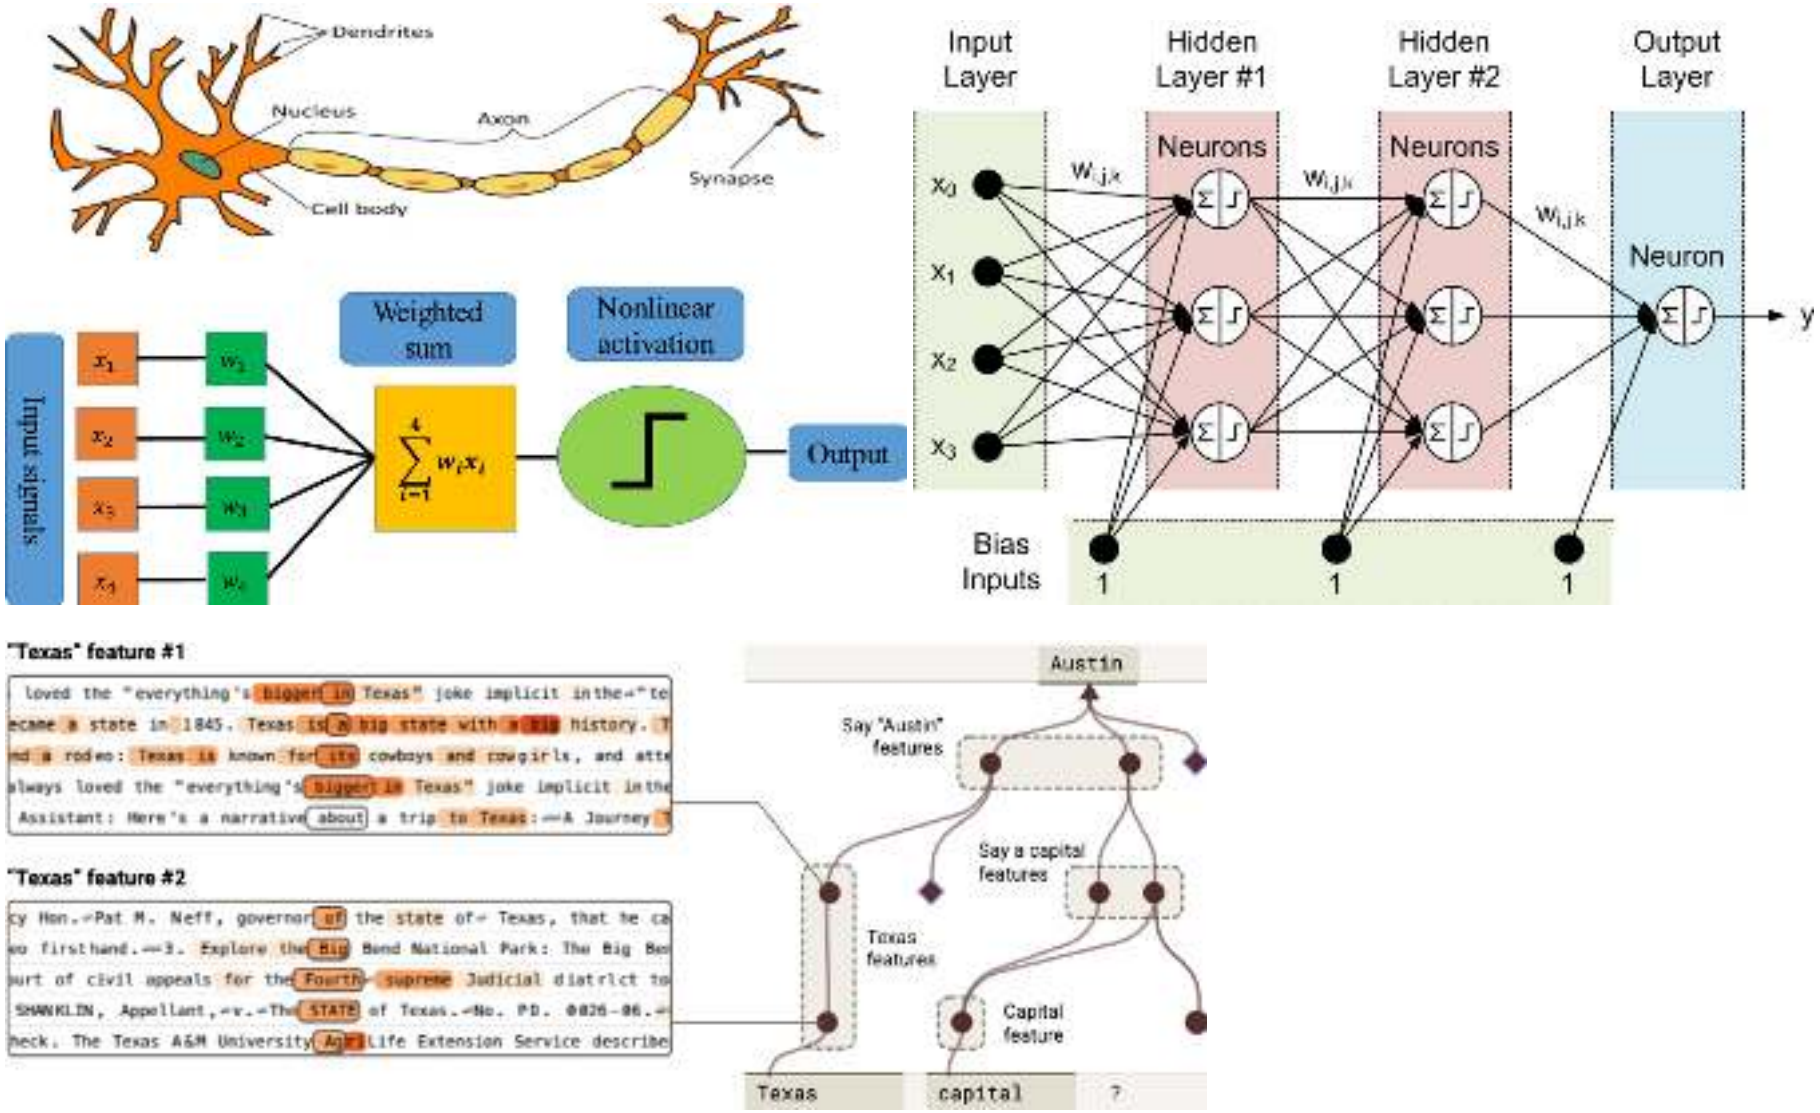

Ameisen, et al., "Circuit Tracing: Revealing Computational Graphs in Language Models", Transformer Circuits, 2025.

Transformer Architecture of LLMs

- 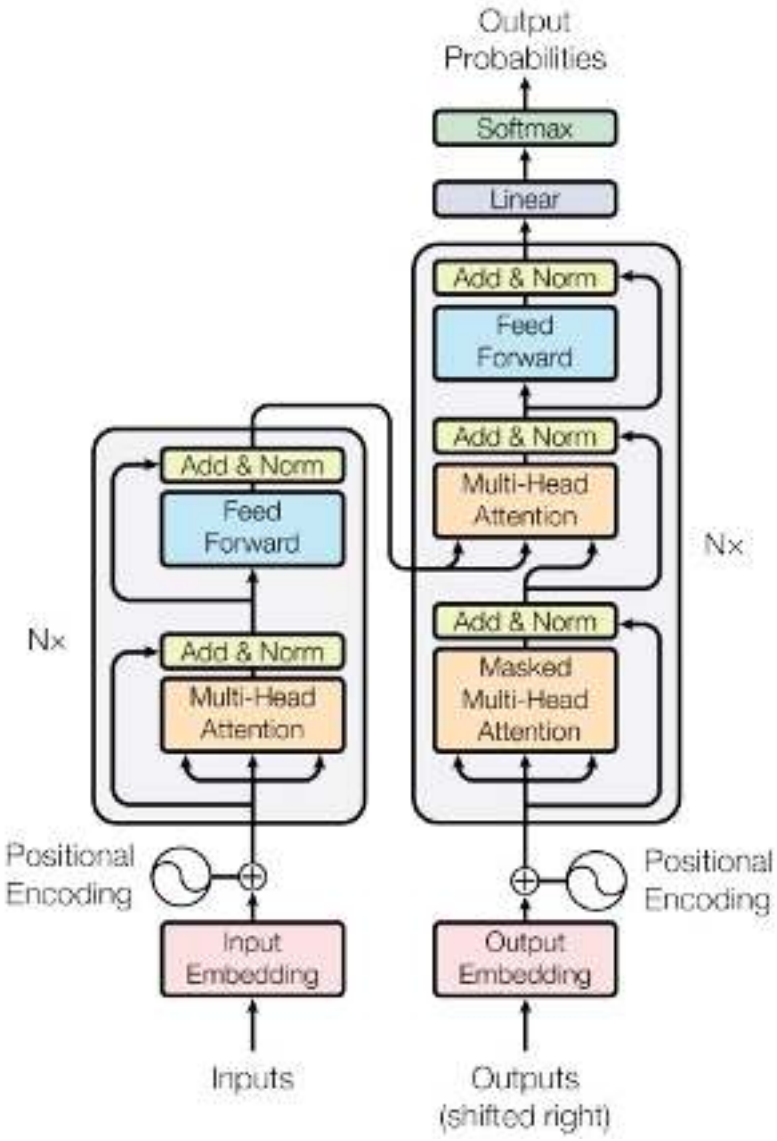

  Attention layer followed by **FF Layer**

- Attention Heads: long-range connections within text

- Feed Forward: same as in previous slide
- N x times : multiple heads run in parallel
- If you trained a neural network, similar plus attention

Vaswani A et al. Attention is all you need. arXiv:1706.03762. 2017;30

The Landscape of Interpretability

- Neurons are "Poly-semantic" (superposition phenomenon)
-

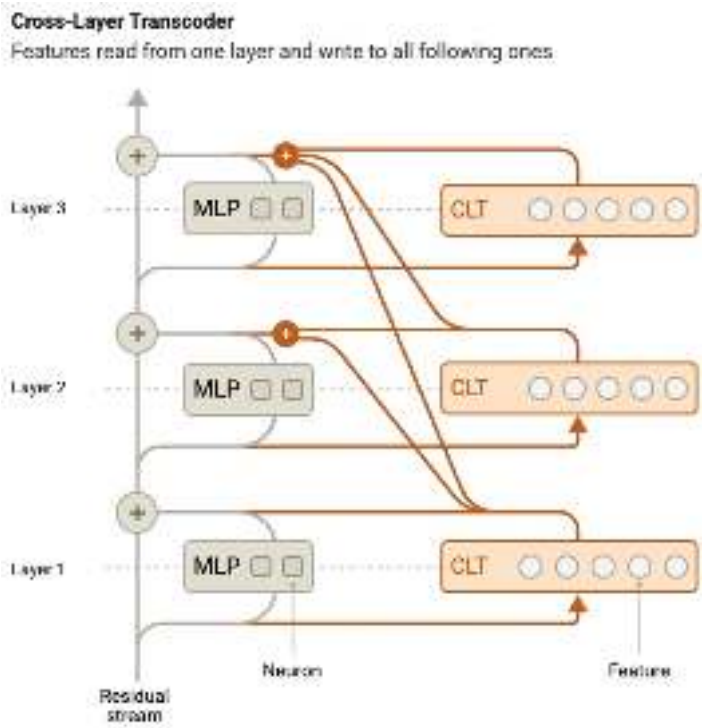

Figure 1: The cross-layer transcoder (CLT) forms the core architecture of our replacement model.

  Sparse Autoencoders: Identifying knowledge (features) stored in the neurons
- Linear Probes: Internal linear representations of specific concepts
- Intervention Experiments: Steering, neural activation patching, and ablations

https://transformer-circuits.pub/

Attribution Graphs

- Interactions between features (text from input) activating neurons
- Graphs showing feature-feature interactions on specific prompts

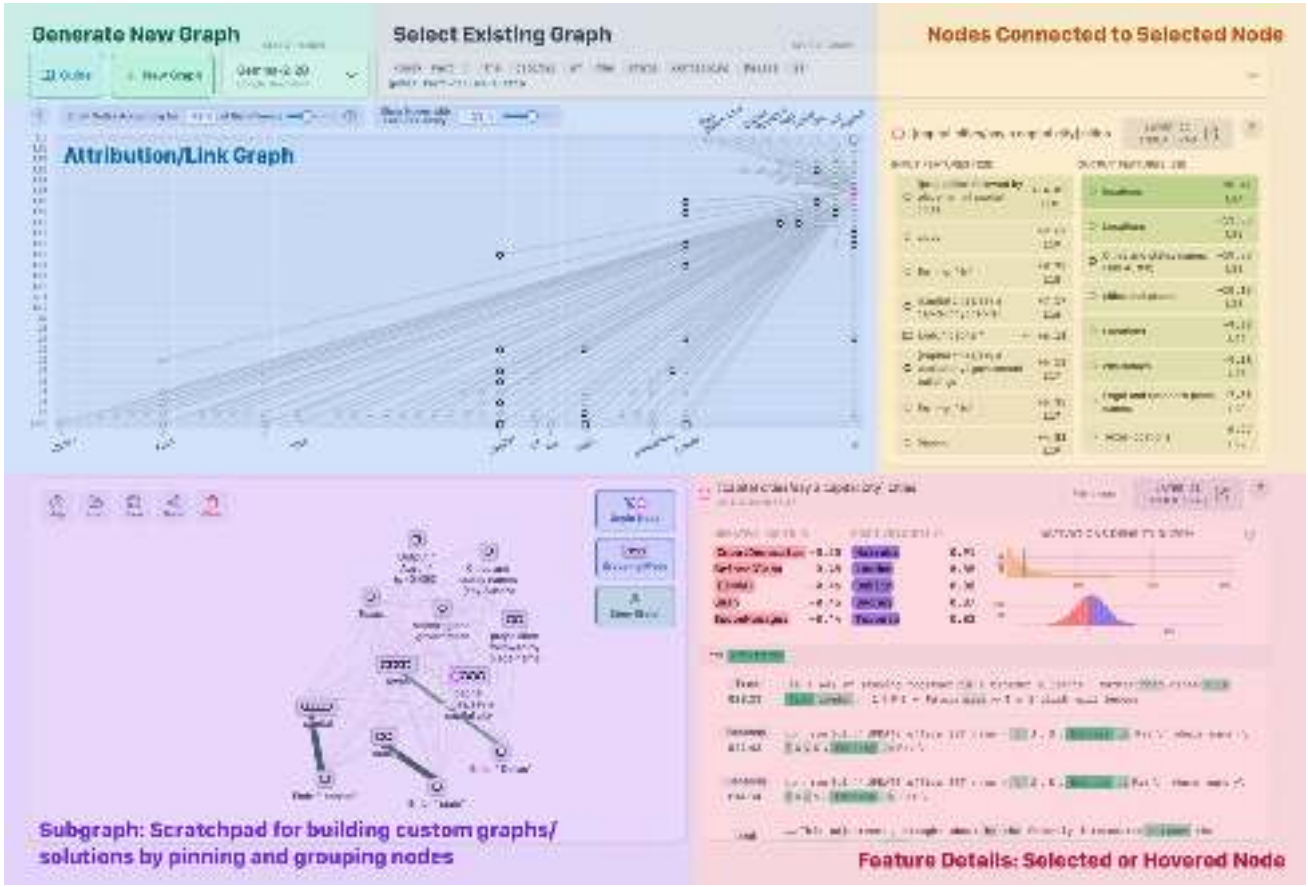

- Interaction chains influencing model output
- Prompt analysis revealing circuits

https://www.neuronpedia.org/

Thanks !

- https://github.com/KKrampis/presentations
- Transformers the tech behind LLMs https://tinyurl.com/3b1b-Transformers
- https://tinyurl.com/alignment-nanda-papers
- https://tinyurl.com/nanda-become-interp-researcher
- https://www.neelnanda.io/mechanistic-interpretability/quickstart